

The Wallfacer Protocol

Containment Without Control:

A Structural Model for Independent AI Oversight

Andrew H. Mangle, Ph.D., CISSP, CISA, CSSLP, PMP

Best Entrust, LLC

Bowie State University

Working Paper

September 2026

Abstract

The rapid advancement of frontier artificial intelligence systems has created a governance gap that existing regulatory frameworks are structurally incapable of closing. Current oversight mechanisms rely on voluntary commitments by AI developers, advisory bodies without binding authority, or government agencies that lack the technical access and speed required to evaluate systems whose capabilities can change between training runs. The result is an industry where the entities with the deepest understanding of the risks are also the entities with the greatest financial incentive to continue development, and where the authority to halt dangerous activity rests almost exclusively with the organizations creating the danger. This paper proposes the Wallfacer Protocol, an independent oversight framework for frontier AI safety modeled on structural principles drawn from nuclear nonproliferation, financial regulation, and judicial independence. The Protocol establishes a small corps of independent overseers (Wallfacers) with the authority to halt AI development activities that pose risks of direct human harm, loss of meaningful human control, or systemic destabilization. Wallfacers operate under comprehensive personal constraints designed to eliminate every identified vector for corruption, capture, or conflicts of interest, including restrictions on financial holdings, political activity, business interests, AI tool usage, media engagement, and web presence. The framework addresses institutional design, funding, selection, authority scope, access mechanisms, knowledge containment, international coordination, and structural resilience. By separating the power to stop from the power to steer, and by placing that power in the hands of individuals whose entire professional existence is structured around independence, the Wallfacer Protocol offers a governance model calibrated to the unique risks of frontier AI development.

I. Introduction

In 2023, leaders from the three largest AI development laboratories signed a joint statement declaring that mitigating the risk of extinction from AI should be a global priority

alongside pandemics and nuclear war. In subsequent months, these same organizations continued to accelerate development timelines, compete for talent and compute resources, and release increasingly capable systems to the public. The dissonance between the stated risk assessment and the observed behavior is not evidence of dishonesty. It is evidence of a structural problem: the entities best positioned to understand AI risk are also the entities least structurally capable of acting on that understanding when action requires slowing down.

This structural problem is not new. It recurs across every domain where the producers of risk are also the primary assessors of that risk. The financial services industry understood the fragility of mortgage-backed securities before the 2008 crisis but could not stop producing them because the competitive pressure to maintain market share overwhelmed the incentive to act prudently. Pharmaceutical companies understood the risks of inadequate clinical trials before regulatory reform, but the commercial pressure to bring drugs to market outpaced internal safety cultures. In both cases, the resolution required the creation of independent bodies with the authority and access to intervene, bodies that were structurally separated from the incentive to continue the risky activity.

The artificial intelligence industry has not yet created such a body. The current landscape of AI governance consists of four categories of oversight, none of which is adequate to the risk profile of frontier systems.

First, voluntary commitments. Several major AI laboratories have published responsible scaling policies, safety frameworks, and deployment guidelines. These commitments are valuable as statements of intent but are unenforceable, revocable, and subject to competitive pressure. When a competitor releases a capable system without equivalent safety measures, the commercial incentive to match that release can override internal safety commitments. No voluntary framework has yet included a binding mechanism for an independent party to halt development.

Second, advisory boards. Companies have established internal and external advisory bodies staffed by ethicists, policy experts, and technical researchers. These boards provide recommendations but lack the authority to compel action. Their influence depends on the

willingness of company leadership to follow advice, which is precisely the dependency that makes self-regulation unreliable.

Third, government agencies. The National Institute of Standards and Technology (NIST) AI Safety Institute, the European Union AI Office, and analogous bodies in other jurisdictions are developing evaluation frameworks and regulatory standards. However, these agencies lack real-time access to frontier model internals, cannot keep pace with development timelines measured in months rather than years, and face the political constraints inherent in government action. Legislation moves slowly; training runs do not.

Fourth, international coordination efforts. The AI Safety Summits, the G7 Hiroshima Process, and bilateral agreements between governments have produced declarations of principle but no operational oversight mechanisms with binding authority over specific development activities.

The Wallfacer Protocol addresses the gap that none of these four categories can fill: the need for independent, technically competent oversight with binding stop authority and real-time access to frontier AI development. The name draws on Liu Cixin's science fiction trilogy *The Three-Body Problem*, in which "Wallfacers" are individuals granted extraordinary authority and resources to defend humanity against an existential threat, operating with a degree of independence that no other institutional role provides. The analogy is deliberate. The Protocol proposes that a small number of individuals, selected for technical competence and ethical grounding and constrained by the most comprehensive set of personal restrictions ever proposed for a governance role, be granted the authority to halt AI development activities that cross defined thresholds of catastrophic risk.

The Protocol's core structural innovation is the separation of access from exploitation. Wallfacers have complete access to frontier model internals, training infrastructure, and capability assessments, but they cannot use what they learn to benefit themselves, their associates, or any commercial entity. They can stop development but cannot redirect it. They can observe but cannot share. Their power is defined entirely by its negative scope: they can say no. They cannot say how.

This paper presents the full structural design of the Wallfacer Protocol, including institutional form, funding mechanism, selection and appointment, the constraint architecture, authority scope and mechanism, access and inspection procedures, knowledge containment, international framework, and structural resilience. It addresses the practical, legal, and philosophical challenges of implementing such a framework and argues that the costs of creating it are justified by the costs of continuing without it.

II. The Case for Independent Oversight with Stop Authority

The argument for independent AI oversight rests on a structural observation, not a moral judgment about the people running AI laboratories. The leaders of frontier AI companies include individuals who are genuinely concerned about safety, who have invested significant resources in alignment research, and who have in some cases publicly acknowledged that their own systems may pose risks they do not fully understand. The problem is not that these individuals are irresponsible. The problem is that responsibility, in the context of competitive markets and fiduciary obligations, is insufficient.

A chief executive of an AI laboratory who concludes that a particular training run should be halted faces a decision matrix in which every incentive pushes toward continuation. Halting development cedes competitive ground to rivals who may not halt. Shareholders and investors expect returns on capital that depend on continued progress. Employees whose compensation includes equity have a financial stake in the company's competitive position. The researcher's own intellectual curiosity and career incentives favor pushing forward. Even the safety team, if it exists, may face internal pressure to approve rather than delay. The result is an environment where stopping is always harder than continuing, even when stopping is the right decision.

This dynamic has a name in the governance literature: regulatory capture by incentive structure. It differs from the more commonly discussed form of regulatory capture, in which a regulator is co-opted by the industry it oversees. Here, the entity that should regulate itself is captured by its own competitive incentives. The distinction matters because it means that replacing the individuals in leadership positions does not solve the

problem. Any competent executive placed in the same structural position will face the same incentive matrix.

The historical evidence for this pattern is extensive. In the period preceding the 2008 financial crisis, risk managers at major banks identified the fragility of collateralized debt obligations and the underpricing of mortgage default risk. Their warnings were overridden not because bank leadership was ignorant but because the competitive pressure to maintain market share and quarterly earnings was structurally stronger than the incentive to act on long-term risk assessments. The resolution required the creation and strengthening of independent regulatory bodies with the authority to impose capital requirements, conduct stress tests, and halt specific activities.

In the pharmaceutical industry, the thalidomide crisis of the late 1950s and early 1960s demonstrated that manufacturers' internal safety assessments were insufficient to protect the public. The result was the empowerment of the Food and Drug Administration to require independent clinical trials and to withhold approval, giving an external body the authority to say no regardless of the manufacturer's assessment of its own product.

In nuclear energy, the Three Mile Island incident in 1979 and the Chernobyl disaster in 1986 exposed the limitations of operator-controlled safety cultures. The response included the creation and strengthening of independent regulatory bodies (the Nuclear Regulatory Commission in the United States, the International Atomic Energy Agency internationally) with the authority to inspect facilities, mandate safety improvements, and shut down reactors that failed to meet standards.

In each case, the pattern is the same. An industry producing systemic risk initially resists independent oversight on the grounds that it understands its own technology better than any external body could. The industry proposes self-regulation as a more efficient and technically informed alternative. Self-regulation functions adequately during periods of stability but fails catastrophically when competitive pressures intensify or when a novel risk emerges that the industry's internal incentive structure is not equipped to address. A crisis occurs. Independent oversight with binding authority is established in the aftermath.

The Wallfacer Protocol proposes that the artificial intelligence industry not wait for its own crisis. The distinguishing characteristic of frontier AI risk, compared to financial, pharmaceutical, or nuclear risk, is irreversibility. A financial crisis destroys wealth but the economy recovers. A pharmaceutical failure harms patients but the damage is bounded. Even a nuclear incident, devastating as it is, affects a limited geographic area. The risk scenarios that frontier AI developers themselves identify as plausible, scenarios involving loss of human control over autonomous systems, weaponization of AI capabilities by state or non-state actors, or systemic disruption of critical infrastructure, are potentially irreversible at a civilizational or species level. If the historical pattern holds and independent oversight is established only after a crisis, the crisis in this domain may not leave a functioning society capable of establishing anything.

III. The Wallfacer Protocol: Institutional Design

A. Institutional Form

The Wallfacer Foundation is conceived as an Independent Oversight Authority: a category of institution that does not map cleanly to existing organizational forms but draws structural elements from several. It is not a government agency, because government agencies are subject to political appointment cycles, congressional funding disputes, and the pace of legislative action, all of which are incompatible with the speed at which frontier AI capabilities advance. It is not a private corporation, because private entities cannot self-grant the coercive authority required to halt another organization's activities. It is not a standard nonprofit, because standard nonprofits lack enforcement mechanisms.

The closest institutional analogs are the Internet Corporation for Assigned Names and Numbers (ICANN), which governs internet domain names under a quasi-governmental charter with international participation; the Federal Reserve System, which operates independently of the executive branch despite being created by statute; and the International Committee of the Red Cross, which holds a unique status under international humanitarian law as a private organization with treaty-based authority and access rights. Each of these institutions demonstrates that it is possible to create a body that exercises

significant authority while maintaining structural independence from both government and industry.

The Wallfacer Foundation follows government policies and operates within the legal frameworks of the jurisdictions in which it functions, but it is not part of any government. Its charter, once established by statute, provides for operational independence comparable to that of a central bank. Its leadership is not politically appointed. Its funding does not flow through appropriations processes. Its decisions are not subject to executive override. This independence is not a luxury; it is a structural requirement. An oversight body that can be defunded by a congressional committee or overruled by an executive order is an oversight body that can be neutralized when its decisions become politically inconvenient.

B. Funding Mechanism

The Foundation is funded by a mandatory levy on frontier AI development activity, indexed to compute expenditure measured in floating-point operations (FLOPs) above a defined threshold. This threshold is set to capture entities training frontier models while excluding small-scale AI users, startups using API access, and academic researchers. The levy functions as a licensing fee: any organization conducting frontier AI development above the compute threshold is required to participate in the oversight framework and fund it proportionally.

The compute-indexed model is preferable to a revenue-based model because compute is a more direct proxy for capability risk. A company might generate modest revenue while training a highly capable and potentially dangerous model; a revenue-based levy would undercount the risk. Compute expenditure, measured in FLOPs, scales directly with the size and capability of the models being trained.

The levy is administered by the Foundation, not by any regulated entity. Companies submit compute usage reports on a quarterly basis, subject to independent audit. The levy rate is set by the Foundation's governance board, subject to review by the same judicial body that handles appeals of stop decisions. This structure parallels the Nuclear Regulatory Commission's funding model, in which licensee fees fund the agency's operations

proportionally to the risk profile of each licensee's activities, and the Federal Deposit Insurance Corporation's insurance premium model, in which bank assessments fund deposit insurance proportionally to each institution's risk profile.

IV. The Wallfacer: Selection, Constraints, and Separation

A. Candidate Profile

A Wallfacer candidate must be technically literate in artificial intelligence systems and information security, ethically grounded through demonstrated professional practice and personal conduct, capable of obtaining the highest level of security clearance in their jurisdiction, and free of current or recent professional, financial, or personal ties to any AI development company. The candidate must also demonstrate a willingness to accept comprehensive personal restrictions that will govern every aspect of their professional and financial life for a minimum of fifteen years.

The scarcity of individuals who meet all of these criteria is a feature of the design, not a flaw. The Wallfacer Protocol does not require hundreds of overseers. It requires a small number of individuals, initially nine, with the rare combination of technical competence, ethical commitment, personal discipline, and willingness to sacrifice professional freedom for the duration of their service. The difficulty of finding such individuals is itself a safeguard: it ensures that the role is filled by people who understand and accept the weight of what they are undertaking.

The requirement that candidates have no ties to AI vendors deserves emphasis. In most regulatory frameworks, recusal protocols manage conflicts of interest case by case. The Wallfacer Protocol takes a different approach: it eliminates the conflicts at the point of selection rather than managing them after appointment. A candidate who has worked for, consulted for, invested in, or maintained significant professional relationships with any AI development company within a defined lookback period (recommended: five years) is ineligible. This restriction narrows the candidate pool significantly but ensures that every

Wallfacer begins service with no entanglement that could compromise their independence or provide grounds for attacking their credibility.

B. Selection Pipeline

The selection process is designed to evolve over the life of the institution. The founding cohort of nine Wallfacers will, by necessity, be volunteers: individuals who step forward from the eligible pool and accept the constraints before the institution has established its reputation, its operational norms, or its political legitimacy. The founding cohort's willingness to serve under conditions of maximal uncertainty is both a practical necessity and a self-selection mechanism. Individuals who volunteer for a role with no precedent, extraordinary personal restrictions, and uncertain institutional support are, almost by definition, motivated by mission rather than self-interest.

As the institution matures, the selection pipeline formalizes. Candidates are nominated by a rotating panel drawn from three pools: academic institutions with relevant programs in artificial intelligence, cybersecurity, and technology governance; national academies of science and engineering; and international standards bodies such as ISO and IEEE. Nominated candidates undergo an independent background investigation equivalent to the process for the highest level of security clearance in their jurisdiction. This investigation examines financial history, foreign contacts, substance use, susceptibility to coercion, and any factor that might compromise the candidate's independence or reliability.

Vetted candidates are confirmed by a supermajority vote of the Foundation's governance board. The board's composition rules prevent any single company, country, or institutional category from holding a blocking vote. Appointments are for non-renewable ten-year terms. The non-renewable feature is essential: it removes the incentive to please any constituency in hopes of reappointment, which is the single most common corruption vector in regulatory appointments.

C. Terms of Service and Separation

A Wallfacer serves a ten-year term, followed by a five-year paid separation period. During the separation period, the former Wallfacer retains all behavioral and financial restrictions but loses stop authority. They continue to receive their full salary. They cannot seek employment in the AI industry, use AI tools, make political contributions, maintain a public web presence, or engage in any activity prohibited during active service.

The separation period serves a specific containment function. A Wallfacer who has spent ten years inspecting frontier AI systems possesses knowledge about model architectures, capabilities, vulnerabilities, and developmental trajectories that would be extraordinarily valuable to any AI company, government intelligence agency, or hostile actor. The five-year separation ensures that this knowledge decays naturally through obsolescence before the former Wallfacer re-enters unrestricted life. In a field where model architectures and capabilities evolve on timescales of months, five years of separation effectively guarantees that the specific technical knowledge a Wallfacer acquired during service will be outdated by the time the restrictions lift.

The total fifteen-year commitment (ten years of active service plus five years of separation) is the longest restriction period proposed for any civilian governance role. It is justified by the nature of the knowledge involved. Unlike financial regulators, whose knowledge of market mechanics remains broadly applicable after leaving service, a Wallfacer's knowledge includes specific details about AI capabilities that could be weaponized or commercially exploited if released prematurely. The restriction period is the simplest and most reliable containment mechanism: rather than attempting to classify and declassify specific pieces of knowledge, the Protocol waits for the knowledge to become irrelevant through the passage of time.

D. The Constraint Architecture

The behavioral and financial restrictions on Wallfacers are not guidelines or best practices. They are contractual conditions of appointment, violation of which constitutes grounds for removal. Each constraint is designed to close a specific, identified vector for corruption, capture, or conflicts of interest.

Compensation. The Wallfacer salary is fixed at one million dollars annually, indexed to inflation and pegged above the median total compensation of senior researchers at frontier AI laboratories. The purpose of the salary level is not to attract candidates through wealth; it is to make the constraints survivable. An individual who cannot own a business, cannot consult, cannot hold individual investments, and cannot seek outside income for fifteen years requires compensation sufficient to maintain financial stability for themselves and their dependents throughout the entire period of service and separation.

Investment restrictions. Wallfacers may hold only broad market index funds (total market or S&P 500 equivalent). Individual stocks, sector funds, cryptocurrency, private investments, and any financial instrument tied to the performance of specific companies or technology sectors are prohibited. This restriction eliminates the most direct financial conflict: a Wallfacer who holds shares in NVIDIA, Microsoft, Google, or any AI-adjacent company has a financial interest in the continued growth of the AI industry that is incompatible with the authority to halt development. Broad market funds allow the Wallfacer to participate in general economic growth without differential benefit from AI-specific outcomes. Full financial disclosure is required annually, reviewed by the Foundation's inspector general.

No AI tool usage. Wallfacers are prohibited from using generative AI, machine learning-based analysis tools, and AI-powered decision support systems in both their professional evaluations and their personal lives. This restriction is not Luddism; it is cognitive independence. A Wallfacer who relies on an AI system to analyze whether AI systems are safe has a dependency that compromises the independence of their judgment. The prohibition draws a clear line: if a tool uses machine learning to generate, predict, or recommend, it is prohibited. Traditional software (spreadsheets, databases, statistical packages, word processors) remains available. The definition of AI tools is subject to review and update on a five-year cycle by the Foundation's governance board, ensuring the restriction remains meaningful as AI capabilities become increasingly embedded in everyday technology.

No business interests. Wallfacers may not own, operate, or hold equity in any business. They may not consult, contract, serve on advisory boards, or accept speaking fees. Their sole source of income is the Foundation salary. This restriction closes the revolving-door problem before it opens. The risk is not limited to direct conflicts of interest; it extends to the appearance of future benefit. When a Wallfacer halts a training run, the affected company's first strategic response will be to attack the Wallfacer's credibility and motives. The absence of any business interest makes that attack structurally impossible to sustain.

No political activity. Wallfacers may not make political contributions, issue political endorsements, engage in partisan activity, or provide testimony before legislative bodies except under subpoena. The AI policy landscape is heavily shaped by campaign contributions and lobbying. A Wallfacer who donates to a senator on the Commerce Committee creates a basis for framing every subsequent stop decision as politically motivated. The restriction extends beyond campaign contributions to encompass any form of political engagement, creating a condition of enforced apoliticism that exceeds the Hatch Act restrictions on federal employees.

No web presence and low public profile. Wallfacers do not maintain personal websites, social media accounts, or any form of public digital presence. They do not publish articles, give interviews, appear at conferences, or participate in public discourse outside of official Foundation communications. Their identity is publicly known (it must be, for the legal authority of their decisions to hold), but they do not cultivate visibility. The Foundation provides a security detail appropriate to the threat level, because a known individual with the authority to halt multi-billion-dollar programs and the knowledge of classified model capabilities is a target for state actors, corporate adversaries, and ideological actors from multiple directions.

Device and interaction monitoring. Wallfacers use Foundation-issued devices for all professional work. These devices run Foundation-controlled operating systems with full activity logging. Personal devices are permitted but must be registered and are subject to periodic audit to verify compliance with the no-AI-tools restriction and communication protocols. All meetings with representatives of regulated companies are logged and

recorded. This monitoring regime is analogous to the random drug testing required of airline pilots and nuclear facility operators: in safety-critical roles, invasive monitoring is a condition of the authority the role carries. The Wallfacer consents to this monitoring as a condition of appointment, and the monitoring itself serves as armor. A Wallfacer whose every professional action is documented cannot be credibly accused of hidden motives or undisclosed conflicts.

Security clearance. Wallfacers must hold the highest level of security clearance available in their jurisdiction. The clearance process serves as an additional vetting mechanism, examining financial history, foreign contacts, and susceptibility to coercion. It also provides a legal framework for prosecuting unauthorized disclosure: rather than creating new criminal penalties for leaking model capabilities information, the existing statutes governing classified information apply.

E. Removal and Graceful Failure

The Wallfacer Protocol treats removal as a system feature, not a punishment. A Wallfacer who fails, whether through judgment error, personal conduct issues, burnout, or simply being wrong about a stop decision, is removed from authority but remains inside the containment structure. They continue to receive their full salary. They continue to be bound by all behavioral and financial restrictions. The ten-year service clock and five-year separation clock apply from their removal date.

Removal falls into three categories. Voluntary resignation: the Wallfacer chooses to leave, but the containment restrictions follow them through the full remaining period. Performance-based removal: the panel or inspector general determines that the Wallfacer is no longer effective, not through misconduct but through declining capability or judgment. Misconduct removal: violation of the constraint architecture triggers investigation and may result in criminal referral depending on the nature of the violation.

In all three categories, the financial and behavioral restrictions continue through the full separation period. The only variable is whether the removal triggers additional investigation. This design eliminates a perverse incentive that plagues other regulatory

roles: a struggling Wallfacer who fears that removal means financial ruin and public disgrace has an incentive to conceal problems and continue operating beyond their capacity. By guaranteeing continued compensation and treating removal as an institutional function rather than a personal failure, the Protocol encourages Wallfacers to self-report concerns about their own fitness before those concerns become operational risks.

V. Authority: Scope, Mechanism, and Limits

A. Trigger Criteria

The Wallfacer's stop authority is activated by three categories of risk, each defined with sufficient specificity to be actionable while retaining the flexibility to address emergent capabilities that were not anticipated at the time of the Protocol's adoption.

Category One: Demonstrated capacity for direct harm. The AI system can be used to synthesize biological or chemical weapons, execute autonomous cyberattacks on critical infrastructure, cause physical harm without human intervention, or enable specific acts of violence or mass casualty events. This category requires demonstrated capability, not theoretical possibility: the system must have shown, through testing or observation, that it can perform the harmful action, not merely that it might be able to in principle.

Category Two: Loss of meaningful human control. The AI system exhibits goal-directed behavior that resists correction, modifies its own training process or objective function without authorization, demonstrates capabilities that were not intended or predicted by its developers and that create new attack surfaces, or operates in ways that its developers cannot explain or replicate. This category captures the scenario that AI developers themselves identify as their primary concern: building a system whose behavior they cannot predict, understand, or override.

Category Three: Systemic destabilization risk. The AI system's deployment would create irreversible changes to critical social infrastructure, including financial markets, electoral systems, military command and control, energy grids, or communications networks, in ways that cannot be recalled or corrected after the fact. This is the broadest

category and carries the highest evidentiary standard: a stop decision under Category Three requires a supermajority of the Wallfacer panel and a detailed technical finding demonstrating the irreversibility of the anticipated impact.

B. Explicit Exclusions

The Wallfacer Protocol is not a mechanism for regulating the pace of technological progress. The following are explicitly excluded from the scope of Wallfacer authority: competitive advantage gained through AI capabilities, market disruption caused by AI-enabled products or services, job displacement resulting from automation, cultural or social changes associated with AI adoption, and any concern that falls under the general category of "this technology makes me uncomfortable." The Wallfacer is a circuit breaker for catastrophic risk, not a gatekeeper for innovation. This exclusion list is as important as the trigger criteria. Without it, the Protocol risks being perceived (and potentially used) as a tool for incumbents to slow competitors or for ideological actors to impose their preferences on technological development. The narrow scope of the stop authority is what makes it politically viable and institutionally legitimate.

C. Tiered Authority

The stop authority operates on three tiers, designed to provide proportional response while maintaining the speed required for genuine emergencies.

Temporary hold. A single Wallfacer can issue an immediate 72-hour halt on a specific training run, deployment, or capability release based on a preliminary assessment that the activity meets one or more trigger criteria. This is the emergency brake. It does not require consensus, committee review, or advance notice. One Wallfacer, present at a facility and observing a concern, can pull it. The 72-hour window provides time for the full panel to be convened and briefed without allowing the potentially dangerous activity to continue.

Extended suspension. Converting a temporary hold into a six-month suspension requires a majority vote of the assigned review panel (three Wallfacers). During the suspension, the company must provide full technical access for the panel's review. The suspension period is

used for thorough evaluation, including independent testing of the system's capabilities against the trigger criteria.

Permanent prohibition. A permanent prohibition on a specific development activity requires a supermajority of the full Wallfacer panel and a written technical finding filed with the sealed record system. The company can appeal a permanent prohibition to an independent arbitration body staffed by technically competent adjudicators, not a general court, because general courts lack the technical expertise to evaluate frontier AI capability claims.

A fourth mechanism addresses emergency deployments. If a company needs to deploy a safety patch or capability restriction urgently (for example, in response to an active exploitation of a vulnerability), requiring even a single Wallfacer review could create dangerous delay. Under the emergency deployment protocol, the company may deploy first and submit to Wallfacer review within 72 hours. The reviewing Wallfacer retains the authority to order a rollback if the review reveals concerns. This mechanism parallels the Food and Drug Administration's Emergency Use Authorization process, in which deployment is permitted under urgent conditions but remains subject to full review.

D. Consequence-Free Escalation

Any single Wallfacer who is uncertain about whether an observed activity meets the trigger criteria can escalate the concern to a full tribunal without that escalation counting against them in any performance assessment, credibility evaluation, or review process. This provision is modeled on aviation safety reporting culture, in which pilots can report near-miss incidents without penalty, a design that has been credited with significantly improving aviation safety outcomes. The Wallfacer Protocol applies the same principle: a system that penalizes uncertainty will produce Wallfacers who suppress uncertainty, which is the most dangerous possible outcome for an oversight body.

VI. Access and Inspection

A. The On-Site Inspection Model

The Wallfacer's access to frontier AI development is modeled on the inspection regime established under the Microsoft antitrust consent decree, adapted for the specific characteristics of AI systems. Under the consent decree, third-party reviewers were granted physical access to Microsoft's source code and development infrastructure in a monitored environment. Reviewers could inspect code and documentation, take limited handwritten notes, but could not bring electronic devices into the inspection area or remove materials from the facility.

Translated to the AI oversight context, this means that Wallfacers have physical access to the laboratory's training infrastructure, model weights, alignment testing results, internal safety evaluations, capability assessments, and source code. They can observe training runs in progress. They can review evaluation results and internal safety reports. They can interview researchers and engineers working on the systems under review. They conduct these inspections on-site, at the laboratory's facilities, using Foundation-issued devices that remain at the facility. Handwritten notes are classified and filed with the Foundation's sealed record system. The laboratory cannot refuse access. The laboratory cannot sanitize or curate what the Wallfacer sees. Laboratory personnel may be present during inspections, consistent with the Protocol's principle that the regulated entities can observe the oversight process.

B. Panel Structure and Rotation

Major new model releases, defined as new foundation models, significant capability jumps, or new modality introductions, require review by a three-person Wallfacer panel. Each member of the panel conducts independent inspections and files an independent assessment. The panel deliberates and reaches a consensus or majority decision on whether to approve the release, impose conditions, or exercise stop authority.

Iterative updates to existing models, including fine-tuning, reinforcement learning from human feedback updates, capability patches, and safety modifications, require review by a single Wallfacer. The single reviewer can escalate to a full panel at any point without consequence if concerns arise during the review.

Wallfacers rotate across laboratory assignments to prevent any single individual from becoming the permanent overseer of any one company. Rotations overlap to ensure that institutional knowledge about a specific laboratory's systems transfers between reviewers. At no point is a single Wallfacer the only person who has inspected a particular system.

C. Whistleblower Protections

The Foundation operates a secure, anonymous reporting channel through which employees of regulated companies can submit concerns about AI safety directly to the oversight body. Reports are received by the Foundation's institutional intake process, not by individual Wallfacers. They are triaged by the inspector general and routed to the appropriate review panel based on the laboratory and system involved. The whistleblower does not know which Wallfacer receives their information, and the Wallfacer does not know who submitted the report unless the whistleblower voluntarily identifies themselves.

This design prevents the development of personal relationships between laboratory employees and specific Wallfacers that could be exploited by either party. It also ensures that whistleblower information enters the institutional record and is evaluated through the standard review process rather than being held by a single individual whose judgment might be affected by the personal dynamics of the reporting relationship.

VII. Knowledge Containment and Accountability

A. The Containment Philosophy

The Wallfacer Protocol inverts the traditional relationship between transparency and secrecy in oversight roles. In most governance frameworks, the overseer's personal life is private and the overseer's findings are public. The Wallfacer Protocol reverses this: the Wallfacer's personal life, financial transactions, communications, and professional activities are transparent (to both the Foundation's inspector general and, through structured read-only access, to the regulated companies), while the knowledge the Wallfacer acquires through inspections is classified.

This inversion serves both practical and philosophical purposes. Practically, the transparency of the Wallfacer's life eliminates the surface area for attacks on their credibility. When a stop decision costs a company billions of dollars, the company's first response will be to investigate the Wallfacer for hidden motives, undisclosed conflicts, or personal biases. A Wallfacer whose every financial transaction and professional interaction is already documented has no hidden vulnerabilities to exploit. The transparency is armor.

Philosophically, the inversion reflects the Protocol's theory of legitimate authority. The Wallfacer's authority derives not from their personal reputation or public standing but from the institutional constraints that guarantee their independence. The public does not need to trust the individual Wallfacer; it needs to trust the structure that selected them, the restrictions that constrain them, and the monitoring that verifies their compliance. The person is transparent. The knowledge is contained. The authority flows from the structure, not the individual.

This philosophical distinction connects directly to the literary source of the Protocol's name. In Liu Cixin's novels, the Wallfacer's power came from their opacity: their plans worked because no one could see inside their decision-making process. The Wallfacer Protocol inverts this. The Wallfacer's life is visible. The containment is on the information, not the person. This is a meaningful philosophical departure from the fictional concept, and it is deliberate: in a governance framework designed for the real world, the legitimacy of extraordinary authority depends on the accountability of the person exercising it.

B. Sealed Record System

Every stop decision is accompanied by a classified technical finding filed with a designated judicial body. The finding details the specific observations, technical analysis, and risk assessment that supported the decision. The finding is sealed for twenty-five years from the date of filing and then declassified and made available to the public.

The twenty-five-year period serves multiple functions. It protects the competitive interests of the company whose activities were halted by preventing competitors from accessing technical details about the company's development program. It protects national security

by preventing the dissemination of information about AI capabilities that could be weaponized. It protects the Wallfacer from real-time reputational damage that might result from controversial decisions. And it provides eventual public accountability, ensuring that every exercise of the stop authority will ultimately be subject to public scrutiny and historical evaluation.

C. Company Challenge Process

A company subject to a stop decision can challenge the decision through closed proceedings before the independent arbitration body. The company has access to the technical finding relevant to its own systems but cannot access findings related to other companies' systems. The arbitration body consists of technically competent adjudicators, not generalist judges, appointed through a process independent of both the Foundation and the regulated industry.

D. Oversight of the Wallfacers

The regulated companies have structured, read-only monitoring access to verify Wallfacer compliance with the behavioral restrictions. Companies can observe that the Wallfacer's financial transactions, communications metadata, and professional activities are consistent with the constraint architecture. They cannot use this monitored information for litigation, lobbying, or public attacks on the Wallfacer personally. Violation of this boundary triggers penalties against the company.

The Foundation's inspector general provides independent oversight of the Wallfacer corps, conducting periodic audits of compliance with the constraint architecture and investigating allegations of misconduct. The inspector general reports to the Foundation's governance board, not to the Wallfacer corps, ensuring independence.

Twice yearly, the full Wallfacer corps convenes at a secure location for calibration, knowledge sharing, and collective assessment of the threat landscape. No outside attendees are permitted. Proceedings are classified and filed with the sealed record system. These retreats serve as both operational calibration sessions and mental health checkpoints: the

Wallfacer role is isolating by design, and regular contact with the only other people who understand the work is a resilience mechanism, not a luxury.

VIII. Structural Resilience

A. Scaling and Cohort Design

The Wallfacer corps begins with a founding cohort of nine volunteers and scales through a rolling onboarding process designed to match the growth of the frontier AI landscape. Three new Wallfacers are onboarded every three years. During the first decade, this produces a growth phase: nine at founding, twelve in year three, fifteen in year six, eighteen at peak capacity in year nine. When the founding cohort completes its ten-year term and enters separation in year ten, the corps returns to nine active members. It rebuilds to twelve in year twelve as the next cohort is onboarded, and the cycle stabilizes into a steady-state oscillation between nine and twelve active Wallfacers.

This scaling model serves several purposes. The growth phase during the first decade provides increasing capacity as the number of regulated entities and the inspection workload grow. The steady-state oscillation between nine and twelve provides natural surge capacity during periods of higher staffing (twelve allows four simultaneous three-person panels) and maintains operational viability at the baseline (nine allows three simultaneous panels or two panels with three single reviewers available for iterative updates). The regular turnover prevents institutional calcification, ensuring that fresh perspectives cycle into the corps every three years while the mentorship structure maintains continuity.

The founding cohort occupies a unique position in the institution's history. They have no mentors, no precedent, and no established operational culture. Their task includes not only conducting the first inspections but also building the inspection protocols, reporting templates, escalation norms, and inter-laboratory comparison frameworks that will guide every subsequent cohort. The sealed record system preserves the founding cohort's

decisions and reasoning as institutional case law, accessible to future Wallfacers as a reference base.

B. Mentorship and Knowledge Transfer

New Wallfacers are onboarded in cohorts of three, staggered every three years. Each new cohort is mentored by the cohort immediately ahead of them during the overlap period. Over a ten-year term, a Wallfacer is mentored for their first three years, operates independently during the middle period, and mentors the incoming cohort during their final three years. This structure ensures that at any given time, the corps includes experienced mentors, independent operators, and developing members simultaneously.

The mentorship structure also provides resilience against the loss of individual members. If a Wallfacer is removed, resigns, or becomes incapacitated, the cohort structure absorbs the loss without creating a single point of failure. Panel rotation ensures that multiple Wallfacers have inspected each regulated system, so no single individual holds unique knowledge about any laboratory's capabilities.

C. Structural Review Cycle

Every five years, the Foundation's governance board conducts a comprehensive structural review. The review addresses the definition of AI tools (to ensure the restriction remains meaningful as AI capabilities evolve), the trigger criteria for stop authority (to incorporate new categories of risk), the compensation level (to maintain its function as a constraint-enabler), the constraint architecture (to identify any gaps or obsolescence), and the operational protocols (to incorporate lessons learned from the preceding five years of inspections). The review is conducted by the governance board with input from the active Wallfacer corps, but the Wallfacers do not have voting authority over changes to the framework. They are consulted, not empowered to modify their own constraints.

IX. International Framework

The international dimension of AI oversight presents the same fundamental challenge that nuclear nonproliferation faced in the mid-twentieth century: the most dangerous technology is being developed by geopolitical competitors who have limited trust in each other and no existing institutional framework for mutual oversight.

The Wallfacer Protocol proposes a phased approach modeled on the history of the Nuclear Non-Proliferation Treaty (NPT) and the International Atomic Energy Agency (IAEA), adapted for the characteristics of AI development.

In the near term, the Foundation operates as a US-based institution covering US-based laboratories and any international laboratory that voluntarily participates. Some international laboratories may choose to participate for reputational value, particularly those based in allied nations with strong technology governance cultures. Voluntary participation establishes the operational norms and demonstrates the viability of the oversight model.

In the medium term, bilateral agreements between allied nations establish reciprocal Wallfacer arrangements with shared standards but nationally managed personnel security. Each participating nation maintains its own Wallfacer corps, selected and constrained according to the Protocol's standards but operating within the legal and security frameworks of their own jurisdiction. A coordination secretariat within the Foundation facilitates information sharing (within the constraints of the sealed record system) and ensures consistency of inspection standards.

In the long term, the bilateral network evolves into a treaty-based model following the NPT/IAEA template. Each signatory nation commits to subjecting its frontier AI development activities to Wallfacer oversight and to funding the international framework proportionally. An international secretariat coordinates inspection assignments, manages cross-border knowledge containment, and facilitates the dispute resolution process.

This paper acknowledges that adversarial nations, most notably China and Russia, are unlikely to participate in the initial phases of the framework. The Wallfacer Protocol does not depend on universal participation for its effectiveness. A US-led coalition of willing

participants covering the majority of global frontier AI development establishes the institutional precedent and creates normative pressure for broader adoption. The NPT itself required decades to achieve near-universal participation, and the IAEA's authority grew incrementally through demonstrated competence and institutional legitimacy rather than through a single founding moment of universal buy-in.

X. The Philosophical Foundation

The Wallfacer Protocol rests on a philosophical claim that deserves explicit articulation: the most legitimate form of oversight is authority that can only say no.

Most governance frameworks grant overseers both the power to prohibit and the power to direct. Financial regulators can block mergers and also set capital requirements.

Environmental agencies can halt pollution and also mandate remediation. Drug regulators can deny approval and also require additional trials. In each case, the overseer's authority includes both negative power (stopping something) and positive power (requiring something). This combination creates opportunities for the overseer to shape the industry in ways that may reflect the overseer's preferences, biases, or interests rather than the public good. The more an overseer can direct, the more valuable it becomes to capture the overseer.

The Wallfacer Protocol strips the positive power entirely. The Wallfacer can halt development. The Wallfacer cannot redirect it. The Wallfacer cannot require a company to build differently, adopt a particular alignment technique, or pursue a specific research direction. The Wallfacer cannot share what they learn with competitors, regulators, or the public. The authority is purely negative: this proceeds, or this does not proceed.

This limitation is the source of the authority's legitimacy. Because the Wallfacer cannot benefit from the exercise of their power (they cannot redirect development toward approaches they favor, share information with entities they support, or shape the industry's trajectory in any direction), the power is self-limiting. The only rational use of the stop authority is to prevent harm. There is no other use it can be put to.

The philosophical framework also addresses the personal cost of the role. The Wallfacer Protocol asks a small number of individuals to bear extraordinary restrictions for an extended period in service of a function that, if it operates correctly, will be invisible. A Wallfacer whose stop decisions prevent catastrophe will never receive public credit, because the catastrophe that was prevented will, by definition, never happen. The sealed record system ensures that the reasoning behind their decisions will not be publicly available during their lifetime or for years afterward. They sacrifice professional freedom, financial flexibility, political participation, public identity, and personal privacy in exchange for the knowledge that they are serving a function that they believe matters.

This is not a role that can be motivated by external rewards. The compensation is adequate, but the restrictions make it impossible to leverage the role into wealth, fame, or influence. The Wallfacer serves because they believe the work is necessary, because they possess the rare combination of competence and willingness that the role requires, and because they accept that the measure of their success is the absence of the disaster they were appointed to prevent.

XI. Conclusion

The frontier AI industry is currently governed by a patchwork of voluntary commitments, advisory bodies, and nascent regulatory frameworks, none of which provides independent, technically competent oversight with binding authority to halt dangerous development. The structural incentives facing AI developers, competitive pressure, fiduciary obligations, and the asymmetry between the short-term costs of stopping and the long-term costs of continuing, make self-regulation inadequate to the risk profile of the technology.

The Wallfacer Protocol offers a structural solution. It creates a small corps of independent overseers with the access, authority, and independence required to serve as a credible check on frontier AI development. It constrains those overseers through the most comprehensive set of personal restrictions ever proposed for a governance role, eliminating every identified vector for corruption, capture, and conflicts of interest. It contains the knowledge those overseers acquire through a sealed record system that

provides eventual public accountability without real-time disclosure of sensitive capabilities information. It scales through a cohort structure that matches the growth of the industry. And it provides a template for international coordination modeled on the most successful nonproliferation framework in history.

The cost of implementing the Protocol is not trivial. It requires legislation, international negotiation, and the identification of individuals willing to accept extraordinary personal sacrifice. But the cost of inaction is clear. The AI industry is producing systems of increasing capability on accelerating timelines, with oversight mechanisms that depend on the willingness of the entities creating the risk to regulate themselves. History demonstrates, across every comparable domain, that this arrangement fails when the stakes are highest.

The question is not whether independent oversight of frontier AI development is necessary. The question is whether it will be established before or after the crisis that makes its necessity undeniable. The Wallfacer Protocol proposes that we not wait to find out.

References

Atomic Energy Act of 1954, Pub. L. No. 83-703, 68 Stat. 919 (1954).

Center for AI Safety. (2023). *Statement on AI risk*. Retrieved from <https://www.safe.ai/statement-on-ai-risk>

Energy Reorganization Act of 1974, Pub. L. No. 93-438, 88 Stat. 1233 (1974) (establishing the Nuclear Regulatory Commission).

European Parliament and Council. (2024). Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). *Official Journal of the European Union*, L 2024/1689.

Federal Deposit Insurance Act, 12 U.S.C. § 1811 et seq. (1950).

Federal Food, Drug, and Cosmetic Act, 21 U.S.C. § 360bbb-3 (Emergency Use Authorization authority, Section 564).

- Financial Crisis Inquiry Commission. (2011). *The financial crisis inquiry report: Final report of the National Commission on the Causes of the Financial and Economic Crisis in the United States*. Washington, DC: U.S. Government Printing Office.
- Hatch Act of 1939, 5 U.S.C. §§ 7321-7326 (restrictions on political activity of federal employees).
- International Atomic Energy Agency. (1956). *Statute of the International Atomic Energy Agency*. Vienna: IAEA.
- Internet Corporation for Assigned Names and Numbers. (1998). *Articles of incorporation and bylaws*. Los Angeles, CA: ICANN.
- Kefauver-Harris Drug Amendments of 1962, Pub. L. No. 87-781, 76 Stat. 780 (1962) (strengthening FDA premarket approval requirements following the thalidomide crisis).
- Kemeny, J. G., et al. (1979). *Report of the President's Commission on the Accident at Three Mile Island*. Washington, DC: U.S. Government Printing Office.
- Liu, C. (2015). *The dark forest* (J. Martinsen, Trans.). New York, NY: Tor Books. (Original work published 2008).
- National Institute of Standards and Technology. (2023). *Artificial intelligence risk management framework (AI RMF 1.0)*. Gaithersburg, MD: U.S. Department of Commerce.
- Sarbanes-Oxley Act of 2002, Pub. L. No. 107-204, 116 Stat. 745 (2002) (establishing the Public Company Accounting Oversight Board).
- Stop Trading on Congressional Knowledge (STOCK) Act of 2012, Pub. L. No. 112-105, 126 Stat. 291 (2012).
- Treaty on the Non-Proliferation of Nuclear Weapons, opened for signature July 1, 1968, 21 U.S.T. 483, 729 U.N.T.S. 161.
- United States v. Microsoft Corp., Final Judgment, Civil Action No. 98-1232 (D.D.C. Nov. 12, 2002) (consent decree establishing third-party technical review of protocols and source code).
- Wall Street Reform and Consumer Protection Act (Dodd-Frank Act) of 2010, Pub. L. No. 111-203, 124 Stat. 1376 (2010).